

Machine Learning HW7

Ben Zhang, bz957

May 11, 2018

3 Ridge Regression

3.1

```
class L2NormPenaltyNode(object):
    """ Node computing  $l2\_reg * ||w||^2$  for scalars  $l2\_reg$  and vector  $w$  """
    def __init__(self, l2_reg, w, node_name):
        """
        Parameters:
        l2_reg: a scalar value  $\geq 0$  (not a node)
        w: a node for which w.out is a numpy vector
        node_name: node's name (a string)
        """
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.l2_reg = l2_reg
        self.w = w

    def forward(self):
        self.out = np.sum(self.w.out ** 2) * self.l2_reg  ###
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        d_w = self.d_out * 2 * self.w.out * self.l2_reg
        self.w.d_out += d_w
        return self.d_out

    def get_predecessors(self):
        return [self.w]
```

3.2

```
class SumNode(object):
    """ Node computing  $a + b$ , for numpy arrays  $a$  and  $b$  """
    def __init__(self, a, b, node_name):
        """
        Parameters:
        a: node for which  $a.out$  is a numpy array
        b: node for which  $b.out$  is a numpy array of the same shape as  $a$ 
        node_name: node's name (a string)
        """
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.a = a
        self.b = b

    def forward(self):
        self.out = self.a.out + self.b.out
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        d_a = self.d_out
        d_b = self.d_out
        self.a.d_out += d_a
        self.b.d_out += d_b
        return self.d_out

    def get_predecessors(self):
        return [self.a, self.b]
```

3.3

```

class RidgeRegression(BaseEstimator, RegressorMixin):
    """ Ridge regression with computation graph """
    def __init__(self, l2_reg=1, step_size=.005, max_num_epochs = 5000):
        self.max_num_epochs = max_num_epochs
        self.step_size = step_size
        # Build computation graph
        self.x = nodes.ValueNode(node_name="x") # to hold a vector input
        self.y = nodes.ValueNode(node_name="y") # to hold a scalar response
        self.w = nodes.ValueNode(node_name="w") # to hold the parameter vector
        self.b = nodes.ValueNode(node_name="b") # to hold the bias parameter (scalar)
        self.l2_reg = nodes.ValueNode(node_name="l2_reg") # to hold the reg parameter
        self.prediction = nodes.VectorScalarAffineNode(x=self.x, w=self.w, b=self.b,
                                                         node_name="prediction")
        self.squareloss = nodes.SquaredL2DistanceNode(a=self.prediction, b=self.y,
                                                         node_name="square_loss")
        self.l2penalty = nodes.L2NormPenaltyNode(l2_reg=l2_reg, w=self.w,
                                                  node_name="l2_penalty")
        self.objective = nodes.SumNode(a=self.squareloss, b=self.l2penalty,
                                       node_name="objective")

        # Group nodes into types to construct computation graph function
        self.inputs = [self.x]
        self.outcomes = [self.y]
        self.parameters = [self.w, self.b]

        self.graph = graph.ComputationGraphFunction(self.inputs, self.outcomes,
                                                    self.parameters,
                                                    self.prediction,
                                                    self.objective)

        # TODO

```

This code passed the test in `ridge_regression.t.py`

```

/Users/zhangben/PycharmProjects/venv/bin/python /Applications/PyCharm.app/Contents/Helpers/RunPythonTools.py --run /Users/zhangben/Google Drive/NYU/Classes/1003MachineLearning/hw/hw7-backprop/c
import sys; print('Python_%s_on_%s' % (sys.version, sys.platform))
sys.path.extend(['/Users/zhangben/Google Drive/NYU/Classes/1003MachineLearning/hw/hw7-backprop/c
DEBUG: (Node l2 norm node) Max rel error for partial deriv w.r.t. w is 1.0008354029219
.DEBUG: (Node sum node) Max rel error for partial deriv w.r.t. a is 1.6365788421260423
DEBUG: (Node sum node) Max rel error for partial deriv w.r.t. b is 1.6365788421260423
.DEBUG: (Parameter w) Max rel error for partial deriv 3.4803029014850967e-09.
DEBUG: (Parameter b) Max rel error for partial deriv 1.0710782025328269e-09.
.

```

Ran 3 tests in 0.003s
OK

For this parameter setting,

```
l2reg = 1
```

```
estimator = RidgeRegression(l2_reg=l2reg, step_size=0.00005, max_num_epochs=2000)
```

The average square error on the training set is about 0.21 (for all the epoch), the average square error of last epoch is 0.19

For this parameter setting,

```
l2reg = 0
```

```
estimator = RidgeRegression(l2_reg=l2reg, step_size=0.0005, max_num_epochs=500)
```

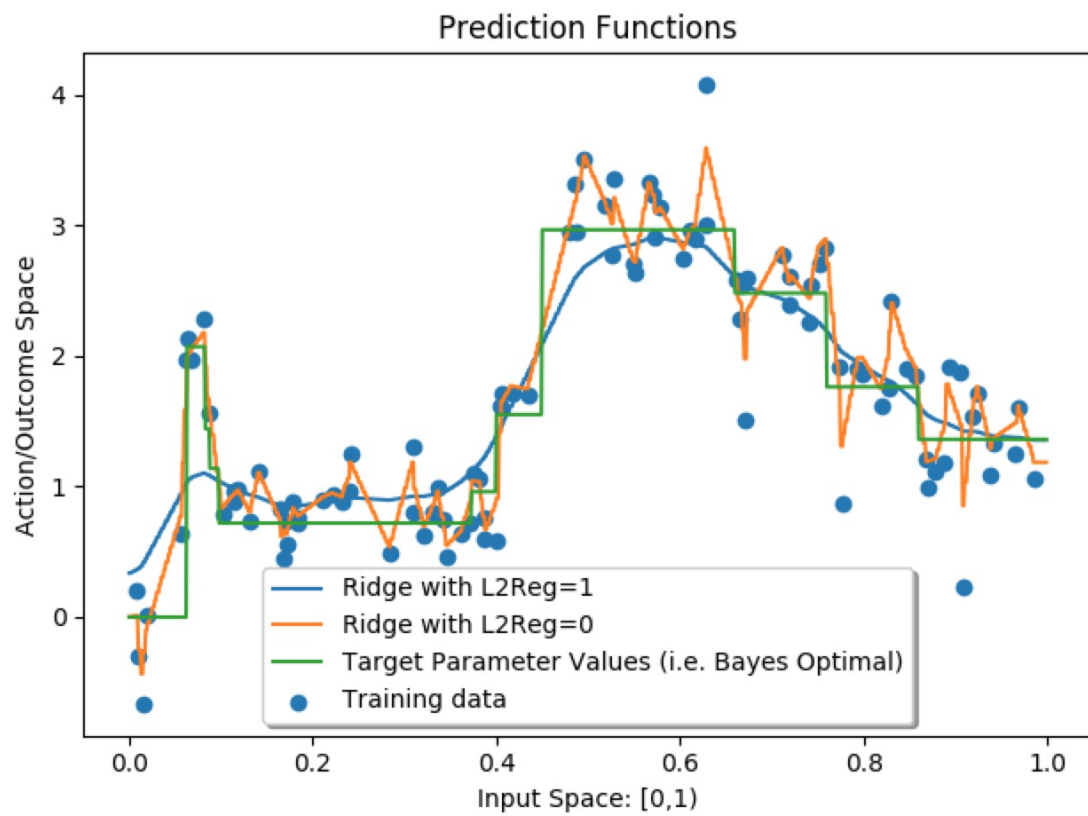
```
estimator.fit(X_train, y_train)
```

```
name = "Ridge_with_L2Reg="+str(l2reg)
```

```
pred_fns.append({ "name": name, "preds": estimator.predict(X) })
```

The average square error on the training set is about 0.09 (for all the epoch), the average square error of last epoch is 0.026

```
/Users/zhangben/PycharmProjects/venv/bin/python "/Users/zhangben/Google Drive/NYU/Classes/1003MachineLearn
Epoch 0 : Ave objective= 1.1824952961993138 Ave training loss: 0.6463960990296029
Epoch 50 : Ave objective= 0.3212158047986074 Ave training loss: 0.21159871027610827
Epoch 100 : Ave objective= 0.31468140111720166 Ave training loss: 0.20004363361489313
Epoch 150 : Ave objective= 0.3165904315213173 Ave training loss: 0.19830474564021358
Epoch 200 : Ave objective= 0.3160686787534271 Ave training loss: 0.1981266965157925
Epoch 250 : Ave objective= 0.31534423038536125 Ave training loss: 0.19683668417796646
Epoch 300 : Ave objective= 0.31614445968323185 Ave training loss: 0.1973942179156053
Epoch 350 : Ave objective= 0.31406617847386803 Ave training loss: 0.19729557981178755
Epoch 400 : Ave objective= 0.3135215863506393 Ave training loss: 0.19893910160053657
Epoch 450 : Ave objective= 0.3142910132707722 Ave training loss: 0.20009184211031544
Epoch 500 : Ave objective= 0.31444083445655996 Ave training loss: 0.19757106438917962
Epoch 550 : Ave objective= 0.3134930817187636 Ave training loss: 0.20010517902913008
Epoch 600 : Ave objective= 0.3138838198746616 Ave training loss: 0.19844564630580777
Epoch 650 : Ave objective= 0.3135285819301993 Ave training loss: 0.1974881801236934
Epoch 700 : Ave objective= 0.3130506662261157 Ave training loss: 0.19732614789517317
Epoch 750 : Ave objective= 0.3117013712459582 Ave training loss: 0.19783101780306273
Epoch 800 : Ave objective= 0.31168129407553313 Ave training loss: 0.1983611483883822
Epoch 850 : Ave objective= 0.31194493882665475 Ave training loss: 0.1975972825831876
Epoch 900 : Ave objective= 0.3091803049094511 Ave training loss: 0.19803753368486882
Epoch 950 : Ave objective= 0.30730845754053315 Ave training loss: 0.20106822952961476
Epoch 1000 : Ave objective= 0.309648352854229 Ave training loss: 0.1988515657685185
Epoch 1050 : Ave objective= 0.30988723244639055 Ave training loss: 0.19805672982522401
Epoch 1100 : Ave objective= 0.3113556507248592 Ave training loss: 0.19895299706976444
Epoch 1150 : Ave objective= 0.31005154841183946 Ave training loss: 0.1979775922437964
Epoch 1200 : Ave objective= 0.3090916231752089 Ave training loss: 0.1995000126381679
Epoch 1250 : Ave objective= 0.3100790723657532 Ave training loss: 0.1986279496123064
Epoch 1300 : Ave objective= 0.3100129202032706 Ave training loss: 0.19810130047319854
Epoch 1350 : Ave objective= 0.30864372505852944 Ave training loss: 0.20018117630983412
Epoch 1400 : Ave objective= 0.30739788393397194 Ave training loss: 0.1991664830783109
Epoch 1450 : Ave objective= 0.30932908915768725 Ave training loss: 0.1986643511216483
Epoch 1500 : Ave objective= 0.31003686035328387 Ave training loss: 0.19950146408211517
Epoch 1550 : Ave objective= 0.3082133897248385 Ave training loss: 0.20080273488480033
Epoch 1600 : Ave objective= 0.30986572689496766 Ave training loss: 0.19871890220213367
Epoch 1650 : Ave objective= 0.3092978425000933 Ave training loss: 0.19843253742878036
Epoch 1700 : Ave objective= 0.30799576359962655 Ave training loss: 0.19977815901456666
Epoch 1750 : Ave objective= 0.3082377093984994 Ave training loss: 0.20104494011928875
Epoch 1800 : Ave objective= 0.30910853213014977 Ave training loss: 0.19924764482104265
Epoch 1850 : Ave objective= 0.3081311104395419 Ave training loss: 0.1990516125737984
Epoch 1900 : Ave objective= 0.3076586631286824 Ave training loss: 0.2016693038929235
Epoch 1950 : Ave objective= 0.30603514269971993 Ave training loss: 0.20955740968713624
Epoch 0 : Ave objective= 0.6199700676565317 Ave training loss: 0.33779698856819307
Epoch 50 : Ave objective= 0.11640087158103882 Ave training loss: 0.14831928241549905
Epoch 100 : Ave objective= 0.09255635105225568 Ave training loss: 0.06511970822795728
Epoch 150 : Ave objective= 0.07341384555517372 Ave training loss: 0.06614559514424374
Epoch 200 : Ave objective= 0.052895571649770455 Ave training loss: 0.07425199570794411
Epoch 250 : Ave objective= 0.05335887862867526 Ave training loss: 0.0451404005220485
Epoch 300 : Ave objective= 0.04680622490251268 Ave training loss: 0.0361065398823835
Epoch 350 : Ave objective= 0.04685768710835255 Ave training loss: 0.03370235147177942
Epoch 400 : Ave objective= 0.043209080599909105 Ave training loss: 0.030626203721320043
Epoch 450 : Ave objective= 0.039207821860192386 Ave training loss: 0.027362762770559015
```



4.1.1 The Affine Transformation

4.1.1

$$1) \frac{\partial J}{\partial w_{ij}} = \sum_{r=1}^m \frac{\partial J}{\partial y_r} \cdot \frac{\partial y_r}{\partial w_{ij}} = \frac{\partial J}{\partial y_1} \cdot \frac{\partial y_1}{\partial w_{ij}} + \dots + \frac{\partial J}{\partial y_m} \cdot \frac{\partial y_m}{\partial w_{ij}}.$$

$$\because y_i = \sum_{j=1}^d (w_{ij} x_j + b).$$

$$\therefore \frac{\partial y_i}{\partial w_{ij}} = \frac{\partial}{\partial w_{ij}} \left(\sum_{j=1}^d w_{ij} x_j + b \right) = x_j \delta_{ij} = x_j$$

$$\text{for } r \neq i, \quad \frac{\partial y_r}{\partial w_{ij}} = 0$$

$$\therefore \frac{\partial J}{\partial w_{ij}} = \frac{\partial J}{\partial y_i} \cdot x_j$$

$$2) \frac{\partial J}{\partial w} \in \mathbb{R}^{m \times d}, \quad \frac{\partial J}{\partial y} \in \mathbb{R}^{m \times 1}, \quad x \in \mathbb{R}^{d \times 1}$$

$$\therefore \frac{\partial J}{\partial w_{ij}} = \frac{\partial J}{\partial y_i} \cdot x_j$$

$$\therefore \frac{\partial J}{\partial w} = \frac{\partial J}{\partial y} \times X^T = \begin{bmatrix} \frac{\partial J}{\partial y_1} x_1 & \dots & \frac{\partial J}{\partial y_1} x_d \\ \vdots & \ddots & \vdots \\ \frac{\partial J}{\partial y_m} x_1 & \dots & \frac{\partial J}{\partial y_m} x_d \end{bmatrix}$$

3). from 1) & 2), we have $\frac{\partial J}{\partial x_i} = \sum_{r=1}^m \frac{\partial J}{\partial y_r} \cdot \frac{\partial y_r}{\partial x_i}$

$$\Rightarrow \frac{\partial J}{\partial x_i} = \sum_{j=1}^m \frac{\partial J}{\partial y_j} w_{ij}$$

$$\Rightarrow \frac{\partial J}{\partial x} = w^T \left(\frac{\partial J}{\partial y} \right)$$

$$4). \frac{\partial J}{\partial b} = \sum_{r=1}^m \frac{\partial J}{\partial y_r} \cdot \frac{\partial y_r}{\partial b} = \sum_{r=1}^m \frac{\partial J}{\partial y_r} = \frac{\partial J}{\partial y}$$

4.1.2 Element-wise Transformers

4.1.2.

$\therefore S$ has the same dimension as A and $\sigma(A)$,

i.e. $\in \mathbb{R}^n$ indexed by a single variable

$\therefore \frac{\partial J}{\partial S}$ and $\sigma'(A) = \frac{\partial S}{\partial A}$ are both $\in \mathbb{R}^n$

$$\therefore \frac{\partial J}{\partial A_i} = \sum_r \frac{\partial J}{\partial S_r} \cdot \frac{\partial S_r}{\partial A_i} = \frac{\partial J}{\partial S_i} \cdot \sigma'(A_i) = \left(\frac{\partial J}{\partial S} \otimes \sigma'(A) \right)_i$$

$$\therefore \frac{\partial J}{\partial A} = \frac{\partial J}{\partial S} \otimes \sigma'(A)$$

4.2 MLP Implementation

4.2.1

```
class AffineNode(object):
```

```
    """Node implementing affine transformation  $(W, x, b) \rightarrow Wx + b$ , where  $W$  is a matrix,  
    and  $x$  and  $b$  are vectors
```

```

    Parameters:
    W: node for which W.out is a numpy array of shape (m,d)
    x: node for which x.out is a numpy array of shape (d)
    b: node for which b.out is a numpy array of shape (m) (i.e. vector of length m)
    """
    def __init__(self, W1, x, b, node_name):
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.W1 = W1
        self.x = x
        self.b = b

    def forward(self):
        self.out = np.dot(self.W1.out, self.x.out) + self.b.out
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        d_W1 = np.outer(self.d_out, self.x.out)
        d_x = np.dot(self.W1.out.T, self.d_out)
        d_b = self.d_out
        self.W1.d_out += d_W1
        self.x.d_out += d_x
        self.b.d_out += d_b
        return self.d_out

    def get_predecessors(self):
        return [self.W1, self.x, self.b]
    ## TODO

```

4.2.2

```

class TanhNode(object):
    """Node tanh(a), where tanh is applied elementwise to the array a
    Parameters:
    a: node for which a.out is a numpy array
    """
    def __init__(self, a, node_name):
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.a = a

    def forward(self):
        self.out = np.tanh(self.a.out)
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        d_a = (1 - np.tanh(self.a.out)**2)*self.d_out
        self.a.d_out += d_a
        return self.d_out

    def get_predecessors(self):
        return [self.a]
    ## TODO

```

4.2.3

```

class MLPRegression(BaseEstimator, RegressorMixin):
    """ MLP regression with computation graph """
    def __init__(self, num_hidden_units=10, step_size=.005, init_param_scale=0.01, max_num_epochs=1000):
        self.num_hidden_units = num_hidden_units
        self.init_param_scale = 0.01
        self.max_num_epochs = max_num_epochs
        self.step_size = step_size

        # Build computation graph
        self.x = nodes.ValueNode(node_name="x") # to hold a vector input
        self.y = nodes.ValueNode(node_name="y") # to hold a scalar response
        self.W1 = nodes.ValueNode(node_name="W1") # to hold the parameter matrix
        self.w2 = nodes.ValueNode(node_name="w2") # to hold the parameter vector
        self.b1 = nodes.ValueNode(node_name="b1") # to hold the bias parameter (vector)
        self.b2 = nodes.ValueNode(node_name="b2") # to hold the bias parameter (scalar)
        self.affine = nodes.AffineNode(W1=self.W1, x=self.x, b=self.b1,
                                       node_name="affine")
        self.tanh = nodes.TanhNode(a=self.affine, node_name="tanh")
        self.prediction = nodes.VectorScalarAffineNode(x=self.tanh, w=self.w2, b=self.b2,
                                                       node_name="prediction")
        self.objective = nodes.SquaredL2DistanceNode(a=self.prediction, b=self.y,
                                                       node_name="square_loss")

        # Group nodes into types to construct computation graph function
        self.inputs = [self.x]
        self.outcomes = [self.y]
        self.parameters = [self.W1, self.b1, self.w2, self.b2]

        self.graph = graph.ComputationGraphFunction(self.inputs, self.outcomes,
                                                    self.parameters, self.prediction,
                                                    self.objective)

        ## TODO

    def fit(self, X, y):
        num_instances, num_fts = X.shape
        y = y.reshape(-1)
        ## TODO: Initialize parameters (small random numbers — not all 0, to break symmetry)

        s = self.init_param_scale
        init_values = {"W1": np.random.standard_normal((self.num_hidden_units,
                                                         num_fts)),
                      "b1": np.random.standard_normal((self.num_hidden_units)),
                      "w2": np.random.standard_normal((self.num_hidden_units)),
                      "b2": np.array(np.random.randn()) }
        self.graph.set_parameters(init_values)

```

This code passed the test in *mlp_regression.t.py*

```
DEBUG: (Node affine) Max rel error for partial deriv w.r.t. W is 1.36372974102e-08.
DEBUG: (Node affine) Max rel error for partial deriv w.r.t. x is 2.17113106792e-09.
DEBUG: (Node affine) Max rel error for partial deriv w.r.t. b is 1.63657896896e-09.
.DEBUG: (Node tanh) Max rel error for partial deriv w.r.t. a is 4.62053093836e-09.
.DEBUG: (Parameter W1) Max rel error for partial deriv 5.50516200215e-07.
DEBUG: (Parameter b1) Max rel error for partial deriv 4.41987841745e-08.
DEBUG: (Parameter w2) Max rel error for partial deriv 1.60618974379e-09.
DEBUG: (Parameter b2) Max rel error for partial deriv 5.83341433218e-10.
.
```

Ran 3 tests in 2.784s

OK

For this parameter setting,

```
estimator = MLPRegression(num_hidden_units=10, step_size=0.001,
                           init_param_scale=.0005, max_num_epochs=5000)
x_train_as_column_vector = x_train.reshape(x_train.shape[0],1) # fit expects a 2-
x_as_column_vector = x.reshape(x.shape[0],1) # fit expects a 2-dim array
```

The average square error on the training set is about 0.26 (for all the epoch), the average square error of last epoch is 0.2

For this parameter setting,

```
estimator = MLPRegression(num_hidden_units=10, step_size=0.0005,
                           init_param_scale=.01, max_num_epochs=500)
```

The average square error on the training set is about 0.27 (for all the epoch), the average square error of last epoch is 0.1

```
Epoch 4350 : Ave objective= 0.21787052489735056 Ave training loss: 0.20907109703833326
Epoch 4400 : Ave objective= 0.21936335849224314 Ave training loss: 0.20873278536023135
Epoch 4450 : Ave objective= 0.21575739984251302 Ave training loss: 0.21247049899088913
Epoch 4500 : Ave objective= 0.2139164061811711 Ave training loss: 0.21209459618823587
Epoch 4550 : Ave objective= 0.21458928018879014 Ave training loss: 0.21147715727708602
Epoch 4600 : Ave objective= 0.21627786902628107 Ave training loss: 0.2078423849872406
Epoch 4650 : Ave objective= 0.21554682177911375 Ave training loss: 0.2083629754779958
Epoch 4700 : Ave objective= 0.21475840474596034 Ave training loss: 0.2083800324209057
Epoch 4750 : Ave objective= 0.2172372542413261 Ave training loss: 0.2071444139441335
Epoch 4800 : Ave objective= 0.21586805162174622 Ave training loss: 0.20793860500505307
Epoch 4850 : Ave objective= 0.21454826316702078 Ave training loss: 0.20673151269863546
Epoch 4900 : Ave objective= 0.2120277483003585 Ave training loss: 0.20942742504061906
Epoch 4950 : Ave objective= 0.21394041497110503 Ave training loss: 0.20756703284005767
Epoch 0 : Ave objective= 2.964237367097052 Ave training loss: 1.9152730770954447
Epoch 50 : Ave objective= 0.14993791042085985 Ave training loss: 0.14548445404902413
Epoch 100 : Ave objective= 0.12704231220538628 Ave training loss: 0.12299736526515767
Epoch 150 : Ave objective= 0.11822350013082848 Ave training loss: 0.11470003747086316
Epoch 200 : Ave objective= 0.11346753480401076 Ave training loss: 0.11038226829966623
Epoch 250 : Ave objective= 0.11033756909596533 Ave training loss: 0.10739508923090915
Epoch 300 : Ave objective= 0.1078779028018294 Ave training loss: 0.10511600404501621
Epoch 350 : Ave objective= 0.10625527177459411 Ave training loss: 0.10337852824848724
Epoch 400 : Ave objective= 0.10458943739804395 Ave training loss: 0.10216547317716809
Epoch 450 : Ave objective= 0.10343198220784604 Ave training loss: 0.10095378967562234
```

Process finished with exit code 0

